

Discussion_1

September 27, 2022

1 ECS 171 Discussion 1

Date: Sep 27, 2022

TA: Dechen Gao (dcgao@ucdavis.edu)

The goal of our discussion:

- Coding of ML algorithms related to homeworks and lectures
 - This week: how to set up environment for your homeworks
 - Next week: linear regression
-

1.1 Outlines

- What is Jupyter?
- Why we need Jupyter?
- Installation: pip & conda
- Colab & Kaggle usage
- Notebook demos
- Export your notebooks

1.2 Jupyter Notebook

1.2.1 Pros & Cons

Interactive (Jupyter) vs. non-interactive.

1. Pros
 - Easy debugging and data exploration.
 - Rich visualization and presentation.
2. Cons
 - Bad version control.
 - Not modular.

1.2.2 Use Jupyter Locally

- **Installation**
 - pip install ([docs](#))
 - conda ([docs](#))
- **Configuration (Optional)**

- Generate and modify configuration files, or use flags when calling Jupyter
 - * e.g., `--notebook-dir`, where to locate your root directory
 - * e.g., `--port`, specify a port number for Jupyter server
- **Client**
 - VS Code: with Jupyter extension that pulls up Jupyter automatically
 - Web browser: you need to pull up Jupyter manually

1.2.3 Use Jupyter Remotely

Remote options

- [Colab](#)
 - Web browser is all you need to access Notebook.
 - File system is a different (talk about this later).
 - CPU/GPU and RAM limited.
- [Kaggle](#)
 - A competition platform that provides notebook/algorithm resources, datasets.
 - Can access Kaggle datasets easily or you can upload and construct your own.
- SSH
 - expose a Jupyter port on server that can be accessed via your SSH client.
 - One option when you need better performance.

Colab File System Here are [file access demos](#) and our [discussion demo](#).

By default you only have temporary file access during each session, and therefore you can only upload again when the session restarts.

However, you can mount Google Drive and upload dataset and Python modules to Google Drive.

```
# Mount Google Drive
from google.colab import drive
drive.mount('/content/gdrive')
# Optional. Change root directory. Use !pwd to verify
import os
os.chdir("/content/gdrive/My Drive/Colab Notebooks/")
# Optional. Append to system paths for interpreter to search for modules
import sys
sys.path.append('/content/gdrive/')
```

Kaggle Demo

- Explore public datasets and manage your own [Kaggle Datasets](#)
- Access datasets in your notebook

1.2.4 Homework Submission (Tentative)

Upload both your Jupyter source and exported PDFs (recommended for HW1) to Gradescope.

- **Required:** source code `.ipynb`
- **Optional:** exported PDF (extra bonus for HW1) plus sources

Why PDF? We grade on Gradescope where PDFs are well supported. We can view comments, and see where points are deducted in PDFs.

Solution 1: Export to PDFs One line command

```
jupyter nbconvert --execute --to pdf your_notebook.ipynb
```

PDF rendering requires [pandoc](#), nbconvert, and LaTeX support.

```
# Install pandoc
choco install pandoc    # For Windows
brew install pandoc     # For MacOS
# Install converting tool
pip install nbconvert   # or conda install
```

Solution 2: If PDF is not working

- Use the following commands to convert iPython file to HTML first

```
jupyter nbconvert --execute --to html your_notebook.ipynb
```

Open HTML with browser and “Print” in PDF format.

- For Jupyter Notebook opened in web browser, select from top bar **File** and **Download as PDF**.

If PDF format is messed up, please first download HTML and then print HTML to PDF in browser.

1.3 GitHub

1.3.1 Client

- Git shell
- GUI client, e.g., GitHub Desktop

1.3.2 Basic Usage

- Create your account and repo
 - Synchronize your local code sources with remote Git repo
 - Create & merge branches (`git branch`)
-

1.4 Markdown Cell Demo

1.5 Heading 1

1.5.1 Heading 2

Heading 3

Heading 4 I'm **bold** text.

I'm *italicized text*.

Blockquotes are here!

Hyperlink: [Canvas Page](#).

```
# Code highlighting
print('Hello, world')
```

Formula:

- $\sum_{x=a}^b f(x)$
 - $\left(\int_a^b f(x) dx \right)$
-

1.5.2 Code Section

```
[1]: print('Hello, world!')
```

Hello, world!

```
[2]: !echo "Hello, world!"
```

"Hello, world!"

1.6 Python Packages

- We will use **pandas** a lot in homeworks to manipulate and analyse datasets.
- For visualization, we have **matplotlib**, **seaborn**, and builtin in packages like **pandas**.

1.6.1 Usage

Use whatever you prefer and look into usage and docs online.

We will also have more demos coming in the next discussions for data manipulation and visualization.

- [pandas docs](#)
- [matplotlib docs](#)
- [seaborn docs](#)

```
[3]: import pandas as pd
data = pd.read_csv('data.csv')
data
```

```
[3]:      cache-misses_1s  node-loads_1s  branch-misses_1s  branch-load-misses_1s  \
0          37514691         2822421         10070509          10559671
1          49901539         3999875         15983066          15905409
2          50400281         3256093         15329054          15610305
3          33857600         2281202         12144856          12142687
```

4	48650176	3510185	16656488	16676601
...
2995	30437560	2326933	12585063	12241054
2996	35121159	2576033	14087715	13595659
2997	34968122	2696992	14058112	13734181
2998	29050123	2382281	11980752	12213550
2999	31203289	2482180	12597138	12576162

	LLC-store-misses_1s	branch-loads_1s	L1-dcache-stores_1s	\
0	4972315	875088331	661114306	
1	7812096	897853137	756436026	
2	7440166	909703981	734120225	
3	6133250	898523534	722133401	
4	6020216	938655692	790604661	
...	
2995	5041214	961290381	807913282	
2996	5959816	993575253	861158899	
2997	4591435	955564749	820924827	
2998	5229134	965928905	788854210	
2999	4960137	1000022539	835397449	

	L1-icache-load-misses_1s	branch-instructions_1s	iTLB-loads_1s	...	\
0	57139716	825641176	1274640	...	
1	73999248	983561267	1721115	...	
2	70138841	948103570	1738119	...	
3	58023478	924673885	1429445	...	
4	69916162	1039331799	1558873	...	
...	
2995	50286201	1012821501	1057812	...	
2996	54518991	1050605365	895227	...	
2997	61481334	1047588424	1282742	...	
2998	62116174	960454788	1293510	...	
2999	53078400	1000095628	963997	...	

	L1-icache-load-misses_5s	branch-instructions_5s	iTLB-loads_5s	\
0	869593	1966467	10449	
1	923435	48598163	20217	
2	454348	1932445	8099	
3	1445012	9134928	10191	
4	1011623	42037945	29691	
...	
2995	4508220	6894462	70212	
2996	5187666	9620870	97990	
2997	2892833	4103961	48325	
2998	17715975	122466994	103075	
2999	4866974	6441573	78765	

	iTLB-load-misses_5s	dTLB-store-misses_5s	dTLB-load-misses_5s	\
0	23199	2277	26126	
1	23926	9173	36467	
2	11133	1097	16018	
3	14334	2200	9918	
4	32148	5706	23413	
...	
2995	98023	9604	111084	
2996	117748	10511	132867	
2997	59845	5794	72812	
2998	249578	86510	321995	
2999	100181	9085	115750	

	dTLB-stores_5s	node-stores_5s	L1-dcache-load-misses_5s	label
0	1212879	16430	261616	1
1	18616934	116077	5774113	1
2	1322034	11231	169059	1
3	5576265	21731	1755962	1
4	25445021	119196	7101670	1
...
2995	5819375	78739	1210270	30
2996	9296473	74737	1528661	30
2997	3867988	43216	728741	30
2998	100615883	1018327	11113398	30
2999	6439933	78603	1215627	30

[3000 rows x 81 columns]

```
[4]: # Select a few columns
data = data[['cache-misses_1s', 'branch-misses_1s', 'LLC-store-misses_1s',
            ↪ 'branch-instructions_1s', 'iTLB-loads_1s', 'label']]
data
```

	cache-misses_1s	branch-misses_1s	LLC-store-misses_1s	\
0	37514691	10070509	4972315	
1	49901539	15983066	7812096	
2	50400281	15329054	7440166	
3	33857600	12144856	6133250	
4	48650176	16656488	6020216	
...	
2995	30437560	12585063	5041214	
2996	35121159	14087715	5959816	
2997	34968122	14058112	4591435	
2998	29050123	11980752	5229134	
2999	31203289	12597138	4960137	

	branch-instructions_1s	iTLB-loads_1s	label
--	------------------------	---------------	-------

```

0          825641176      1274640      1
1          983561267      1721115      1
2          948103570      1738119      1
3          924673885      1429445      1
4         1039331799      1558873      1
...          ...          ...    ...
2995        1012821501      1057812     30
2996        1050605365       895227     30
2997        1047588424      1282742     30
2998         960454788      1293510     30
2999        1000095628       963997     30

```

[3000 rows x 6 columns]

```
[5]: data.shape
```

```
[5]: (3000, 6)
```

```
[6]: # What's in data?
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   cache-misses_1s                       3000 non-null   int64
1   branch-misses_1s                      3000 non-null   int64
2   LLC-store-misses_1s                   3000 non-null   int64
3   branch-instructions_1s                 3000 non-null   int64
4   iTLB-loads_1s                         3000 non-null   int64
5   label                                 3000 non-null   int64
dtypes: int64(6)
memory usage: 140.8 KB

```

```
[7]: data.describe()
```

```

[7]:      cache-misses_1s  branch-misses_1s  LLC-store-misses_1s  \
count      3.000000e+03      3.000000e+03      3.000000e+03
mean      3.574264e+07      1.320227e+07      5.776213e+06
std       7.392098e+06      2.180458e+06      1.223957e+06
min       6.707911e+06      5.306539e+06      9.561360e+05
25%      3.119621e+07      1.185055e+07      4.980937e+06
50%      3.365164e+07      1.256955e+07      5.609685e+06
75%      3.777566e+07      1.393846e+07      6.375148e+06
max       6.644830e+07      2.228168e+07      1.157133e+07

```


	branch-instructions_1s	iTLB-loads_1s	label
count	3.000000e+03	3.000000e+03	3000.000000
mean	9.891446e+08	1.098288e+06	15.500000
std	1.098516e+08	3.104263e+05	8.656884
min	3.090067e+08	2.007580e+05	1.000000
25%	9.139577e+08	8.916890e+05	8.000000
50%	9.727026e+08	1.039418e+06	15.500000
75%	1.044031e+09	1.216392e+06	23.000000
max	1.374149e+09	2.516269e+06	30.000000

```
[8]: import seaborn as sns
from matplotlib import pyplot as plt
# Color style
sns.set_style('whitegrid')
# Set figure size
plt.figure(figsize=(8, 6))
sns.scatterplot(data=data, x='branch-instructions_1s', y='cache-misses_1s')
```

[8]: <matplotlib.axes._subplots.AxesSubplot at 0x24f7fc65cf8>

```
[9]: # Or you can use pandas builtin to plot
data.plot.scatter(x='branch-instructions_1s', y='label')
```

[9]: <matplotlib.axes._subplots.AxesSubplot at 0x24f7ff59f60>





```
[10]: # Pearson correlation by default in pandas
corr = data.corr()
# This is an example using pandas built-in visualization
# Determine the background colors with values in cells
corr.style.set_precision(3).background_gradient(cmap='RdBu', vmin=-1, vmax=1)
```

```
[10]: <pandas.io.formats.style.Styler at 0x24f3cfed860>
```

```
[11]: # Let's use Seaborn to do the heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(corr, vmin=-1, vmax=1, center=0, annot=False, cmap= 'RdBu')
```


```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x24f679666d8>
```



```
[12]: sns.pairplot(data=data,
                # Use kde for the diagonal subplots
                vars=['branch-misses_1s', 'LLC-store-misses_1s',
                    ↪ 'branch-instructions_1s'],
                diag_kind='kde',
                plot_kws=dict(
                    size=.5,
                    alpha=.5,
```

```
))
```

```
[12]: <seaborn.axisgrid.PairGrid at 0x24f3d49cef0>
```



```
[13]: import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

df = pd.DataFrame(data=np.random.normal(0, 1, (20, 10)))

df_normalized=(df - df.mean()) / df.std()
pca = PCA(n_components=df.shape[1])
```

```
pca.fit(df_normalized)

loadings = pd.DataFrame(pca.components_.T,
columns=['PC%s' % _ for _ in range(len(df_normalized.columns))],
index=df.columns)

plt.plot(pca.explained_variance_ratio_)
plt.ylabel('Explained Variance')
plt.xlabel('Components')
plt.show()
```

